

```
In [1]: import numpy as np
import pandas as pd
```

```
In [2]: df=pd.read_csv('https://github.com/philoma/Dataset/raw/main/Bike%20Prices.csv')
```

```
In [3]: df.head
```

```
Out[3]: <bound method NDFrame.head of
rice Year Seller_Type \
0      TVS              TVS XL 100      30000  2017  Individual
1    Bajaj              Bajaj ct 100      18000  2017  Individual
2      Yo              Yo Style      20000  2011  Individual
3    Bajaj      Bajaj Discover 100      25000  2010  Individual
4    Bajaj      Bajaj Discover 100      24999  2012  Individual
...      ...              ...      ...      ...      ...
1056 Royal      Royal Enfield Electra 5 S      90000  2012  Individual
1057 Hero              Hero Honda Hunk      20000  2010  Individual
1058 Bajaj      Bajaj Pulsar 220 DTS-i      60000  2014  Individual
1059 Hero      Hero Honda CBZ extreme      40000  2009  Individual
1060 Hero      Hero Honda CBZ extreme      35000  2012  Individual

      Owner  KM_Driven  Ex_Showroom_Price
0    1st owner      8000      30490.0
1    1st owner     35000      32000.0
2    1st owner     10000      37675.0
3    1st owner     43000      42859.0
4    2nd owner     35000      42859.0
...      ...      ...      ...
1056 1st owner     40000      NaN
1057 1st owner     17000      NaN
1058 1st owner     16000      NaN
1059 1st owner     50000      NaN
1060 1st owner     60000      NaN

[1061 rows x 8 columns]>
```

```
In [59]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1061 entries, 0 to 1060
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Brand                  1061 non-null   object
1   Model                  1061 non-null   object
2   Selling_Price          1061 non-null   int64
3   Year                   1061 non-null   int64
4   Seller_Type            1061 non-null   object
5   Owner                  1061 non-null   object
6   KM_Driven              1061 non-null   int64
7   Ex_Showroom_Price      626 non-null    float64
dtypes: float64(1), int64(3), object(4)
memory usage: 66.4+ KB
```

```
In [60]: df.shape
```

```
Out[60]: (1061, 8)
```

```
In [61]: #df['Ex_Showroom_Price'].fillna(df.groupby(['Brand'])['Selling_Price'].transform(
df=df.dropna())
```

```
In [62]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 626 entries, 0 to 625
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Brand                  626 non-null    object
1   Model                  626 non-null    object
2   Selling_Price          626 non-null    int64
3   Year                   626 non-null    int64
4   Seller_Type            626 non-null    object
5   Owner                  626 non-null    object
6   KM_Driven              626 non-null    int64
7   Ex_Showroom_Price      626 non-null    float64
dtypes: float64(1), int64(3), object(4)
memory usage: 44.0+ KB
```

In [63]: `df.describe()`

Out[63]:

	Selling_Price	Year	KM_Driven	Ex_Showroom_Price
count	626.000000	626.000000	626.000000	6.260000e+02
mean	59445.164537	2014.800319	32671.576677	8.795871e+04
std	59904.350888	3.018885	45479.661039	7.749659e+04
min	6000.000000	2001.000000	380.000000	3.049000e+04
25%	30000.000000	2013.000000	13031.250000	5.485200e+04
50%	45000.000000	2015.000000	25000.000000	7.275250e+04
75%	65000.000000	2017.000000	40000.000000	8.703150e+04
max	760000.000000	2020.000000	585659.000000	1.278000e+06

==

In [64]: `df.columns`

Out[64]: Index(['Brand', 'Model', 'Selling_Price', 'Year', 'Seller_Type', 'Owner', 'KM_Driven', 'Ex_Showroom_Price'], dtype='object')

In [65]: `df[['Brand']].value_counts()`

Out[65]:

Brand	
Honda	170
Bajaj	143
Hero	108
Yamaha	94
Royal	40
TVS	23
Suzuki	18
KTM	6
Mahindra	6
Kawasaki	4
UM	3
Activa	3
Harley	2
Vespa	2
BMW	1
Hyosung	1
Benelli	1
Yo	1

dtype: int64

```
In [66]: df.columns
```

```
Out[66]: Index(['Brand', 'Model', 'Selling_Price', 'Year', 'Seller_Type', 'Owner',  
              'KM_Driven', 'Ex_Showroom_Price'],  
              dtype='object')
```

```
In [67]: df[['Model']].value_counts()
```

```
Out[67]: Model  
Honda Activa [2000-2015]          23  
Honda CB Hornet 160R            22  
Bajaj Pulsar 180                 20  
Yamaha FZ S V 2.0               16  
Bajaj Discover 125              16  
..  
Royal Enfield Thunderbird 500    1  
Royal Enfield Continental GT [2013 - 2018] 1  
Royal Enfield Classic Stealth Black 1  
Royal Enfield Classic Squadron Blue 1  
Yo Style                        1  
Length: 183, dtype: int64
```

```
In [68]: df[['Seller_Type']].value_counts()
```

```
Out[68]: Seller_Type  
Individual    623  
Dealer        3  
dtype: int64
```

```
In [69]: df.replace({'Seller_Type':{'Individual':0,'Dealer':1}},inplace=True)
```

```
/usr/local/lib/python3.7/dist-packages/pandas/core/indexing.py:1773: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy)
```

```
self._setitem_single_column(ilocs[0], value, pi)
```

```
In [70]: df[['Owner']].value_counts()
```

```
Out[70]: Owner  
1st owner    556  
2nd owner    66  
3rd owner     3  
4th owner     1  
dtype: int64
```

```
In [71]: df.replace({'Owner':{'1st owner' : 0,
'2nd owner':1,
'3rd owner': 2,
'4th owner' : 3 }},inplace=True)
```

/usr/local/lib/python3.7/dist-packages/pandas/core/indexing.py:1773: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
self._setitem_single_column(ilocs[0], value, pi)
```

```
In [72]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 626 entries, 0 to 625
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Brand                  626 non-null   object
1   Model                  626 non-null   object
2   Selling_Price          626 non-null   int64
3   Year                   626 non-null   int64
4   Seller_Type            626 non-null   int64
5   Owner                  626 non-null   int64
6   KM_Driven              626 non-null   int64
7   Ex_Showroom_Price      626 non-null   float64
dtypes: float64(1), int64(5), object(2)
memory usage: 44.0+ KB
```

```
In [73]: y=df['Selling_Price']
```

```
In [74]: y.shape
```

```
Out[74]: (626,)
```

```
In [75]: df.columns
```

```
Out[75]: Index(['Brand', 'Model', 'Selling_Price', 'Year', 'Seller_Type', 'Owner',
               'KM_Driven', 'Ex_Showroom_Price'],
              dtype='object')
```

```
In [76]: X=df[['Year', 'Seller_Type', 'Owner',
               'KM_Driven', 'Ex_Showroom_Price']]
```

In [77]: X.shape

Out[77]: (626, 5)

In [78]: X

Out[78]:

	Year	Seller_Type	Owner	KM_Driven	Ex_Showroom_Price
0	2017	0	0	8000	30490.0
1	2017	0	0	35000	32000.0
2	2011	0	0	10000	37675.0
3	2010	0	0	43000	42859.0
4	2012	0	1	35000	42859.0
...
621	2014	0	3	6500	534000.0
622	2011	0	0	12000	589000.0
623	2017	0	1	13600	599000.0
624	2019	0	0	2800	752020.0
625	2013	0	1	12000	1278000.0

626 rows × 5 columns

▬

In [79]: `from sklearn.model_selection import train_test_split`

In [80]: `X_train,X_test,y_train,y_test=train_test_split(X,y,train_size=0.7,random_state=25)`

In [81]: `from sklearn.linear_model import LinearRegression`

In [82]: `lr=LinearRegression()`

In [83]: `lr.fit(X_train,y_train)`

Out[83]: `LinearRegression()`

In [84]: `y_pred=lr.predict(X_test)`

In [85]: `y_pred.shape`

Out[85]: (188,)

In [86]: `y_pred`

Out[86]: `array([27210.52271465, 56340.08335163, 63471.94671996, 53627.63844785,
 55612.75744268, 53888.92259719, 33751.35275102, 60311.4950183 ,
 113713.05684467, 76639.49332954, 27826.7399381 , 49919.83255841,
 65886.64311457, 26755.12664064, 48277.75426038, 127646.56079335,
 70047.10661635, 39350.67963653, 36081.03597878, 45360.79436339,
 48079.89470577, 44803.02464799, 55161.44026111, 71041.51821318,
 91689.22699159, 49301.53594645, 55988.19326252, 108171.54600296,
 32771.06897901, 25468.20072996, 17128.61806164, 179271.41130746,
 45698.99857622, 31371.09285079, 67886.52106737, 41492.49575815,
 56855.22238602, 47820.47003468, 74682.14053958, 24984.21822736,
 55374.00513699, 41412.36775222, 67991.60287764, 26553.59421844,
 89788.69870689, 45764.83633686, 133888.03770389, 106988.113825 ,
 71176.40667714, 25332.25485946, 79512.43778826, 63914.38088173,
 28632.12110986, 53656.13623937, -5396.37132904, 70377.44571174,
 33313.03576476, 53994.92478411, 67509.85836352, 59735.05378847,
 22199.83644217, 15374.18984158, 44510.76819427, 30279.52476752,
 108243.77037514, 19291.8895874 , 53614.312976 , 59230.23269131,
 60174.2108109 , 45924.63468736, 25770.81883496, 63471.36257814,
 242123.45729792, 61387.72544548, 56510.98127074, 48123.28087213,
 51668.27442011, 90279.76190495, 14827.76533556, 112437.70820504,
 35066.88027405, 30902.41069172, 31441.48921433, 125593.75847157,
 27705.38813164, -11590.29205553, 15582.17108685, 75113.64511232,
 504085.44522282, 123545.42050116, 74770.89327697, 50747.47663245,
 44174.3618212 , 25426.7156106 , 30298.3052462 , 47625.67836414,
 27850.37544807, 28845.23330928, 31580.38624692, 32309.63375635,
 47979.16788554, 65955.46375944, 13432.28218017, 15368.80064986,
 31973.23052409, 110353.92870546, 68181.49509136, 23143.49139797,
 53194.65732076, 34603.36376989, 56002.50967868, 62432.66994305,
 391470.77533201, 3558.29480891, 36019.18494305, 70876.34866549,
 72890.00667025, 137596.01384364, 27620.36308877, 135789.30486854,
 39674.40366791, 58367.0924453 , 42401.21202624, 61864.4379567 ,
 42688.89652842, 63710.34571021, 10604.39360071, 38458.82820943,
 112251.84744225, 115403.00577536, 13658.41734785, 36196.83359584,
 54146.22998932, 97297.85724851, 55029.68137265, 22923.26533437,
 104569.97029689, 41965.75852017, 38759.68546491, 28930.61369011,
 45231.66612551, 48475.43422775, 26739.7225731 , 53598.65972203,
 32558.54954524, 32212.22834942, 68172.98738422, 71839.47716461,
 32003.46692215, 40652.69995971, 39935.92211843, 63444.41846202,
 44545.5818771 , 120873.38389616, 60926.58683174, 62641.82167496,
 60816.47379994, 27098.95433573, 26803.64749618, 48956.00468627,
 62032.88118713, 26471.97495723, 104937.23068766, 132903.3578847 ,
 37469.2040942 , 57579.12080094, 40371.00915736, -7039.40662503,
 26485.40030077, 90782.42554145, 52153.21149321, 56453.74542453,
 80440.59426003, 31890.46870273, 49505.97985573, 24288.36959514,
 25540.47481573, 117708.26333955, 23399.66596746, 63678.40865459,
 70144.29372668, 33434.89010059, 60885.29444481, 58389.55370878,
 35118.7040348 , 58729.4540196 , 34627.9532246 , 38583.4623973])`

In [89]: `from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score`

```
In [90]: mean_absolute_error(y_test,y_pred)
```

```
Out[90]: 12225.7370104107
```


```
In [91]: mean_squared_error(y_test,y_pred)
```

```
Out[91]: 554715615.5043668
```

```
In [92]: r2_score(y_test,y_pred)
```

```
Out[92]: 0.8810414402984937
```

```
In [94]: import matplotlib.pyplot as plt
plt.scatter(y_test,y_pred)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title('Actual Price Vs Predicted Price')
plt.show()
```



Get Future Predictions:

Steps: 1. extract a random row using sample function 2. separate X and y 3. predict

```
In [95]: df_new=df.sample(1)
```


In [96]: df_new

Out[96]:

	Brand	Model	Selling_Price	Year	Seller_Type	Owner	KM_Driven	Ex_Showroom_Price
612	Royal	Royal Enfield Classic Stealth Black	160000	2018	0	0	7500	204977.0

□

In [97]: df_new.shape

Out[97]: (1, 8)

In [98]: X_new=df_new.drop(['Brand', 'Model', 'Selling_Price'],axis=1)

In [99]: y_pred_new=lr.predict(X_new)

In [101]: print(y_pred_new)

[150105.70011051]

Thus 150105.70011051 is the predicted price of Royal Enfield Classic Stealth Black.