

Sales Prediction using Random Forest

Import Libraries

```
In [1]:
import pandas as pd
import numpy as np
```

```
In [2]:
df = pd.read_csv(r'https://raw.githubusercontent.com/YBI-Foundation/Dataset/main/Big%20Sale')
```

```
In [3]:
df.head()
```

Out[3]:

	Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_I
0	FDT36	12.3	Low Fat	0.111448	Baking Goods	33.4874	
1	FDT36	12.3	Low Fat	0.111904	Baking Goods	33.9874	
2	FDT36	12.3	LF	0.111728	Baking Goods	33.9874	
3	FDT36	12.3	Low Fat	0.000000	Baking Goods	34.3874	
4	FDP12	9.8	Regular	0.045523	Baking Goods	35.0874	

In [4]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14204 entries, 0 to 14203
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Item_Identifier                       14204 non-null  object
1   Item_Weight                           11815 non-null  float64
2   Item_Fat_Content                       14204 non-null  object
3   Item_Visibility                       14204 non-null  float64
4   Item_Type                             14204 non-null  object
5   Item_MRP                             14204 non-null  float64
6   Outlet_Identifier                     14204 non-null  object
7   Outlet_Establishment_Year            14204 non-null  int64
8   Outlet_Size                           14204 non-null  object
9   Outlet_Location_Type                 14204 non-null  object
10  Outlet_Type                           14204 non-null  object
11  Item_Outlet_Sales                     14204 non-null  float64
dtypes: float64(4), int64(1), object(7)
memory usage: 1.3+ MB
```

Get columns

In [5]:

```
df.columns
```

Out[5]:

```
Index(['Item_Identifier', 'Item_Weight', 'Item_Fat_Content', 'Item_Visibilit
y',
      'Item_Type', 'Item_MRP', 'Outlet_Identifier',
      'Outlet_Establishment_Year', 'Outlet_Size', 'Outlet_Location_Type',
      'Outlet_Type', 'Item_Outlet_Sales'],
      dtype='object')
```

In [6]:

```
df.describe()
```

Out[6]:

	Item_Weight	Item_Visibility	Item_MRP	Outlet_Establishment_Year	Item_Outlet_Sales
count	11815.000000	14204.000000	14204.000000	14204.000000	14204.000000
mean	12.788355	0.065953	141.004977	1997.830681	2185.836320
std	4.654126	0.051459	62.086938	8.371664	1827.479550
min	4.555000	0.000000	31.290000	1985.000000	33.290000
25%	8.710000	0.027036	94.012000	1987.000000	922.135101
50%	12.500000	0.054021	142.247000	1999.000000	1768.287680
75%	16.750000	0.094037	185.855600	2004.000000	2988.110400
max	30.000000	0.328391	266.888400	2009.000000	31224.726950

In [7]:

```
df['Item_Weight'].fillna(df.groupby(['Item_Type'])['Item_Weight'].transform('mean'), inplace=True)
```

In [8]:

```
df.info()
```

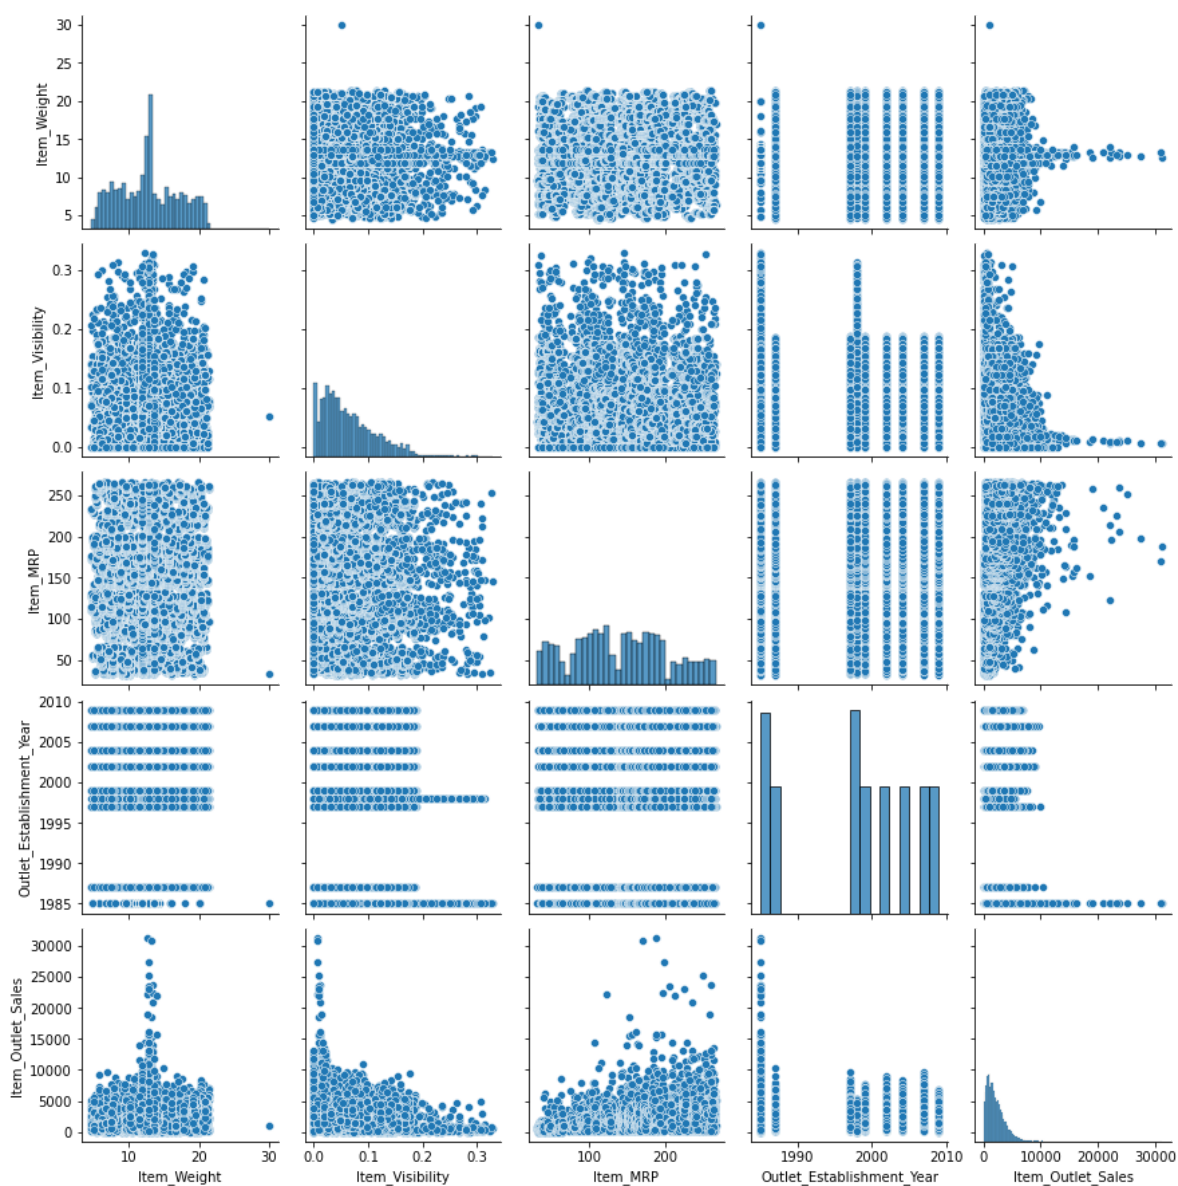
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14204 entries, 0 to 14203
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Item_Identifier                       14204 non-null  object
1   Item_Weight                           14204 non-null  float64
2   Item_Fat_Content                       14204 non-null  object
3   Item_Visibility                       14204 non-null  float64
4   Item_Type                             14204 non-null  object
5   Item_MRP                             14204 non-null  float64
6   Outlet_Identifier                     14204 non-null  object
7   Outlet_Establishment_Year             14204 non-null  int64
8   Outlet_Size                           14204 non-null  object
9   Outlet_Location_Type                  14204 non-null  object
10  Outlet_Type                           14204 non-null  object
11  Item_Outlet_Sales                     14204 non-null  float64
dtypes: float64(4), int64(1), object(7)
memory usage: 1.3+ MB
```

In [9]:

```
import seaborn as sns
sns.pairplot(df)
```

Out[9]:

<seaborn.axisgrid.PairGrid at 0x25601e2af40>



In [10]:

```
df[['Item_Identifier']].value_counts()
```

Out[10]:

```
Item_Identifier
FDQ08          10
FD024          10
FDQ19          10
FDQ28          10
FDQ31          10
..
FDM52           7
FDM50           7
FDL50           7
FDM10           7
FDR51           7
Length: 1559, dtype: int64
```

In [11]:

```
df.replace({'Item_Fat_Content': {'LF': 'Low Fat', 'reg': 'Regular', 'low fat': 'Low Fat'}}, i
```

In [12]:

```
df[['Item_Fat_Content']].value_counts()
```

Out[12]:

```
Item_Fat_Content
Low Fat          9185
Regular          5019
dtype: int64
```

In [13]:

```
df.replace({'Item_Fat_Content': {'Low Fat' : 0, 'Regular' : 1}}, inplace = True)
```

In [14]:

```
df[['Item_Type']].value_counts()
```

Out[14]:

```
Item_Type
Fruits and Vegetables    2013
Snack Foods              1989
Household                1548
Frozen Foods            1426
Dairy                   1136
Baking Goods            1086
Canned                  1084
Health and Hygiene       858
Meat                    736
Soft Drinks              726
Breads                   416
Hard Drinks              362
Others                   280
Starchy Foods           269
Breakfast                186
Seafood                   89
dtype: int64
```

In [15]:

```
cs':0, 'Others' : 2, 'Starchy Foods': 0, 'Breakfast' : 0, 'Seafood' : 0 }}, inplace = True)
```

In [16]:

```
df[['Item_Type']].value_counts()
```

Out[16]:

```
Item_Type
0          11518
1           2406
2            280
dtype: int64
```

In [17]:

```
df[['Outlet_Identifier']].value_counts()
```

Out[17]:

```
Outlet_Identifier
OUT027          1559
OUT013          1553
OUT035          1550
OUT046          1550
OUT049          1550
OUT045          1548
OUT018          1546
OUT017          1543
OUT010           925
OUT019           880
dtype: int64
```

In [18]:

```
'OUT046' : 3, 'OUT035' : 4, 'OUT045' : 5, 'OUT018' : 6, 'OUT017' : 7, 'OUT010' : 8, 'OUT019' :
```

In [19]:

```
df[['Outlet_Identifier']].value_counts()
```

Out[19]:

```
Outlet_Identifier
0                1559
1                1553
2                1550
3                1550
4                1550
5                1548
6                1546
7                1543
8                 925
9                 880
dtype: int64
```

In [20]:

```
df[['Outlet_Size']].value_counts()
```

Out[20]:

```
Outlet_Size
Medium      7122
Small       5529
High        1553
dtype: int64
```

In [21]:

```
df.replace({'Outlet_Size' : {'Small' : 0, 'Medium' : 1, 'High' : 2}}, inplace = True)
```

In [22]:

```
df[['Outlet_Size']].value_counts()
```

Out[22]:

```
Outlet_Size
1                7122
0                5529
2                1553
dtype: int64
```

In [23]:

```
df[['Outlet_Location_Type']].value_counts()
```

Out[23]:

```
Outlet_Location_Type
Tier 3                5583
Tier 2                4641
Tier 1                3980
dtype: int64
```

In [24]:

```
df.replace({'Outlet_Location_Type': {'Tier 1': 0, 'Tier 2': 1, 'Tier 3': 2}}, inplace = True)
```

In [25]:

```
df[['Outlet_Location_Type']].value_counts()
```

Out[25]:

```
Outlet_Location_Type
2                5583
1                4641
0                3980
dtype: int64
```

In [26]:

```
df[['Outlet_Type']].value_counts()
```

Out[26]:

```
Outlet_Type
Supermarket Type1    9294
Grocery Store        1805
Supermarket Type3    1559
Supermarket Type2    1546
dtype: int64
```

In [27]:

```
Outlet_Type' : {'Grocery Store':0, 'Supermarket Type1' : 1, 'Supermarket Type2' : 2, 'Superma
```

In [28]:

```
df[['Outlet_Type']].value_counts()
```

Out[28]:

```
Outlet_Type
1                9294
0                1805
3                1559
2                1546
dtype: int64
```


In [29]:

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14204 entries, 0 to 14203
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Item_Identifier                       14204 non-null  object
1   Item_Weight                           14204 non-null  float64
2   Item_Fat_Content                       14204 non-null  int64
3   Item_Visibility                       14204 non-null  float64
4   Item_Type                             14204 non-null  int64
5   Item_MRP                             14204 non-null  float64
6   Outlet_Identifier                     14204 non-null  int64
7   Outlet_Establishment_Year             14204 non-null  int64
8   Outlet_Size                           14204 non-null  int64
9   Outlet_Location_Type                  14204 non-null  int64
10  Outlet_Type                           14204 non-null  int64
11  Item_Outlet_Sales                     14204 non-null  float64
dtypes: float64(4), int64(7), object(1)
memory usage: 1.3+ MB
```

In [30]:

df.head()

Out[30]:

	Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_I
0	FDT36	12.3	0	0.111448	0	33.4874	
1	FDT36	12.3	0	0.111904	0	33.9874	
2	FDT36	12.3	0	0.111728	0	33.9874	
3	FDT36	12.3	0	0.000000	0	34.3874	
4	FDP12	9.8	1	0.045523	0	35.0874	

Get Shape of Dataframe

In [31]:

df.shape

Out[31]:

(14204, 12)

In [32]:

y = df['Item_Outlet_Sales']

In [33]:

```
y.shape
```

Out[33]:

```
(14204,)
```

In [34]:

```
y
```

Out[34]:

```
0      436.608721
1      443.127721
2      564.598400
3     1719.370000
4      352.874000
...
14199   4984.178800
14200   2885.577200
14201   2885.577200
14202   3803.676434
14203   3644.354765
Name: Item_Outlet_Sales, Length: 14204, dtype: float64
```

In [35]:

```
df.columns
```

Out[35]:

```
Index(['Item_Identifier', 'Item_Weight', 'Item_Fat_Content', 'Item_Visibilit
y',
      'Item_Type', 'Item_MRP', 'Outlet_Identifier',
      'Outlet_Establishment_Year', 'Outlet_Size', 'Outlet_Location_Type',
      'Outlet_Type', 'Item_Outlet_Sales'],
      dtype='object')
```

In [36]:

```
X = [['Item_Identifier', 'Item_Weight', 'Item_Fat_Content', 'Item_Visibility',
      'Item_Type', 'Item_MRP', 'Outlet_Identifier',
      'Outlet_Establishment_Year', 'Outlet_Size', 'Outlet_Location_Type',
      'Outlet_Type']]
```

In [37]:

```
X = df.drop(['Item_Identifier', 'Item_Outlet_Sales'], axis = 1)
```

In [38]:

```
X.shape
```

Out[38]:

```
(14204, 10)
```

In [39]:

X

Out[39]:

	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	C
0	12.300000	0	0.111448	0	33.4874		2
1	12.300000	0	0.111904	0	33.9874		7
2	12.300000	0	0.111728	0	33.9874		6
3	12.300000	0	0.000000	0	34.3874		9
4	9.800000	1	0.045523	0	35.0874		7
...
14199	12.800000	0	0.069606	0	261.9252		4
14200	12.800000	0	0.070013	0	262.8252		7
14201	12.800000	0	0.069561	0	263.0252		1
14202	13.659758	0	0.069282	0	263.5252		0
14203	12.800000	0	0.069727	0	263.6252		2

14204 rows × 10 columns

Get X Variables Standardized

In [41]:

```
from sklearn.preprocessing import StandardScaler
```

In [42]:

```
sc = StandardScaler()
```

In [43]:

```
X_std = df[['Item_Weight', 'Item_Visibility', 'Item_MRP', 'Outlet_Establishment_Year']]
```

In [44]:

```
X_std = sc.fit_transform(X_std)
```

In [45]:

```
X_std
```

Out[45]:

```
array([[ -0.11541705,  0.88413635, -1.73178716,  0.13968068],
       [ -0.11541705,  0.89300616, -1.72373366,  1.09531886],
       [ -0.11541705,  0.88958331, -1.72373366,  1.3342284 ],
       ...,
       [  0.00220132,  0.07011952,  1.96538148, -1.29377659],
       [  0.20444792,  0.06469366,  1.97343499, -1.53268614],
       [  0.00220132,  0.07334891,  1.97504569,  0.13968068]])
```

In [46]:

```
]] = pd.DataFrame(X_std, columns = [['Item_Weight', 'Item_Visibility', 'Item_MRP', 'Outlet_
```

In [47]:

```
X
```

Out[47]:

	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	C
0	-0.115417	0	0.884136	0	-1.731787		2
1	-0.115417	0	0.893006	0	-1.723734		7
2	-0.115417	0	0.889583	0	-1.723734		6
3	-0.115417	0	-1.281712	0	-1.717291		9
4	-0.703509	1	-0.397031	0	-1.706016		7
...
14199	0.002201	0	0.070990	0	1.947664		4
14200	0.002201	0	0.078898	0	1.962160		7
14201	0.002201	0	0.070120	0	1.965381		1
14202	0.204448	0	0.064694	0	1.973435		0
14203	0.002201	0	0.073349	0	1.975046		2

14204 rows × 10 columns

Get train Test Split

In [48]:

```
from sklearn.model_selection import train_test_split
```

In [49]:

```
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size = 0.1, random_state = 25
```

In [51]:

```
X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

Out[51]:

```
((12783, 10), (1421, 10), (12783,), (1421,))
```

Get Model Train

In [52]:

```
from sklearn.ensemble import RandomForestRegressor
```

In [53]:

```
rfr = RandomForestRegressor (random_state = 2539)
```

In [54]:

```
rfr.fit(X_train, y_train)
```

Out[54]:

```
RandomForestRegressor(random_state=2539)
```

Get Model Prediction

In [55]:

```
y_pred = rfr.predict(X_test)
```

In [56]:

```
y_pred.shape
```

Out[56]:

```
(1421,)
```

In [57]:

```
y_pred
```

Out[57]:

```
array([1428.48491758,  739.39517005, 1764.06852049, ..., 2131.8834375 ,  
       3221.29313926,  448.15959596])
```

Get Model Evaluation

In [59]:

```
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
```

In [60]:

```
mean_squared_error(y_test, y_pred)
```

Out[60]:

1617511.846634074

In [61]:

```
mean_absolute_error(y_test, y_pred)
```

Out[61]:

830.7489828870267

In [62]:

```
r2_score(y_test, y_pred)
```


Out[62]:

0.5789856819214596

Get Visualization of Actual vs Predicted Results

In [64]:

```
import matplotlib.pyplot as plt
plt.scatter(y_test, y_pred)
plt.xlabel("Actual Prices")
plt.ylabel("Predicted Prices")
plt.title("Actual Price vs Predicted Price")
plt.show()
```



In []: